

ArcadeDB: One ACID Engine for Documents, Graphs, Vectors, and Time Series

[TODO: working title; alternatives in .notes strategy doc]

Taewoon Kim
HumemAI
taewoon@humem.ai

Roberto Franchini
Arcade Data
r.franchini@arcadedata.com

Christian Himpe
[TODO: affiliation]
[TODO: email]

Luca Garulli
Arcade Data
l.garulli@arcadedata.com

Abstract—[TODO: Written last. **Spine: multi-model unification as a real, adopted system, one page/WAL/MVCC-transaction pipeline under documents, graphs, key-value, time series, and vectors; every index type (LSM, full-text, geo, hash, dense/sparse vector) commits inside the same transaction; Raft replication ships model-agnostic WAL page diffs; one jar runs embedded in-process or as a server speaking six wire protocols. Deep dive: what adding the vector model to a unified substrate takes, including the deliberate transactional-source-of-truth / eventually-consistent-ANN-graph tradeoff. Evaluation vs specialist engines + lessons from a decade of multi-model engineering (OrientDB lineage).]**

Index Terms—multi-model databases, ACID transactions, vector search, graph databases, embedded databases

I. INTRODUCTION

[TODO: Vector-era motivation: apps hold documents + relationships + embeddings of the same entities; composed specialist stacks pay consistency/ops glue. Thesis: unification at the storage layer, not the API gateway. Claims 1–4 from the audit. Explicitly claim integration, not new algorithms.]

II. BACKGROUND AND DESIGN GOALS

[TODO: PROSE PASS NEEDED. Raw co-author input below (Garulli, 2026-07-08).]

ArcadeDB descends from OrientDB, and its design goals are best understood as deliberate corrections of three OrientDB limitations its authors lived with:

- **Durability.** OrientDB’s write-ahead log was asynchronous, single-threaded, and (in the authors’ assessment) unreliable. ArcadeDB’s journal is synchronous, parallel, and crash-proof. It follows the point-of-no-return discipline detailed in §III.
- **Edge storage.** OrientDB managed edges with a tree-bag structure; ArcadeDB replaces it with a LIFO linked list that scales substantially better at high edge counts.
- **High availability.** The OrientDB multi-master model “simply doesn’t work” (authors’ words); ArcadeDB abandons it for a Raft leader/follower architecture (§VI).

[TODO: Design goals framing; what industry usage demanded. Third-party academic adoption: DatAasee [1] (meta-data catalog for research libraries) and AptBacterialDB [2] (ArcadeDB/openCypher graph layer for an antibacterial-aptamer database), both independent multi-model graph uses.

Christian’s “why we chose it” paragraph goes here or in Lessons.]

III. ARCHITECTURE: ONE SUBSTRATE, FIVE MODELS

[TODO: Page -¿ WAL -¿ optimistic-MVCC transaction pipeline; PagedComponent inventory (buckets, LSM trees, vector segments, time-series buckets); index changes committed in commit1stPhase; RID-linked adjacency + LightEdges (honest framing); time-series engine paragraph.]

IV. CASE STUDY: ADDING THE VECTOR MODEL

[TODO: Dense (JVector HNSW) + LSM_SPARSE_VECTOR (BlockMax-WAND DAAT, RID-delta compression, quantization, size-tiered compaction). The tradeoff, stated plainly: vector records are transactional/WAL’d/Raft-replicated; the ANN graph is an asynchronously maintained, rebuildable derived structure. Include the measurement-integrity story (retracted early number) in a sidebar-style paragraph if space allows.]

V. QUERY AND DEPLOYMENT SURFACES

[TODO: One executor + result model behind SQL/Cypher/GraphQL/Mongo (Gremlin via TinkerPop over the same storage); embedded in-process vs server; six wire protocols; cite the SciPy paper for Python binding internals.]

VI. HIGH AVAILABILITY

[TODO: Ratis-based Raft shipping shared-WAL page diffs =¿ model-agnostic replication; quorum/membership hardening lessons.]

VII. EVALUATION

[TODO: E1 sparse DB-vs-DB (Qdrant, Milvus) at 100k/1M/10M; E2 unified-ACID hybrid transaction vs composed stack incl. atomicity demonstration; E3 durability (kill -9) + Raft failover; E4 embedded vs server deployment axis; E5 dense sanity. Protocol: docker, pinned digests, identical cpuset/memory caps, N=5 mean±std, serial re-runs for all reported numbers.]

VIII. LESSONS LEARNED AND TRADEOFFS

[TODO: The cost of unification (measured); where specialists win; single-node ceiling; async-ANN consequences; production war stories (co-authors).]

IX. RELATED WORK

[TODO: Multi-model (Cosmos DB API-level vs our storage-level; ArangoDB; Lu & Holubová survey; UniBench/M2Bench), vector DBs (Milvus, Manu, Qdrant, pgvector), sparse retrieval (SPLADE, BlockMax-WAND), embedded engines (SQLite, DuckDB).]

X. CONCLUSION

[TODO:]

ACKNOWLEDGMENT

[TODO: AI-generated-content disclosure REQUIRED by ICDE 2027: name tools and affected content areas (drafting/editing assistance, experiment tooling). Keep updated as the paper evolves.]

REFERENCES

- [1] C. Himpe, “DatAasee – a metadata-lake as metadata catalog for a virtual data-lake,” *arXiv preprint arXiv:2409.05512*, 2024.
- [2] N. Bajiya, I. Gupta, and G. P. S. Raghava, “AptBacterialDB: A comprehensive, manually curated database of antibacterial aptamers,” bioRxiv preprint, 2026, doi:10.64898/2026.07.01.735956; uses Arcad-eDB/openCypher for the aptamer–target knowledge graph.